

CSC 555: Mining Big Data

Project, Phase 1 (due Sunday, February 14th)

In this part of the project (which will serve as our take-home midterm), you will 1) Set up a 3-node cluster and 2) perform data warehousing and transformation queries using Hive, Pig and Hadoop streaming on that cluster. The modified Hive-style schema is.

http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/SSBM1/SSBM_schema_hive.sql

(you still have to add the delimiter to table definitions)

It is based on SSBM benchmark (derived from industry standard TPCB benchmark). The data is at Scale1, or the smallest unit – lineorder is the largest table at about 0.6GB. You can use wget to download the following links. Keep in mind that data is | -separated.

<http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/SSBM1/dwdate.tbl>

<http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/SSBM1/lineorder.tbl>

<http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/SSBM1/part.tbl>

<http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/SSBM1/supplier.tbl>

<http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/SSBM1/customer.tbl>

Please be sure to submit all code (pig, python and HiveQL).

Part 1: Multi-node cluster

- 1) Your first step is to setup a multi-node cluster and re-run a simple wordcount. For this part, you will create a 3-node cluster (with a total of 1 master + 2 worker nodes). Include your master node in the “slaves” file, to make sure **all 3** nodes are working.

You need to perform the following steps:

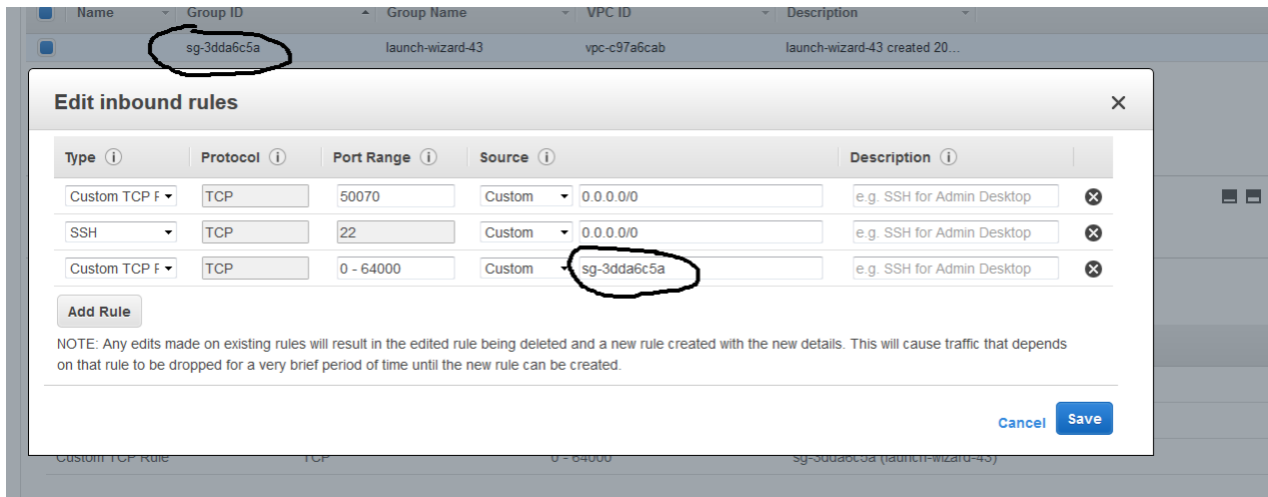
1. Create a new node of a medium size (you can always switch the size of the node). It is possible, but I do not recommend trying to reconfigure your existing Hadoop into this new cluster (it is much easier to make 3 new nodes for a total of 4 in your AWS account).
 - a. **When creating a node I recommend changing the default 8G hard drive to 30G on all nodes.**
 - b. Change your security group setting to open firewall access. We need to open the ports in two different ways. We will open port 50070 for the web interface in order to be able to see the cluster status in a browser. We will also set 0-64000 range opening up all ports. However, we will ensure that the ports are open only **within** the cluster and not to the world.

In order to make changes, you need to do the following. Access the cluster security group (launch-wizard-xx).

Elastic IPs	
Availability zone	us-west-1b
Security groups	launch-wizard-39. view rules
Scheduled events	-

Right click on the security group and choose Edit inbound rules

Note that the first line below is opening port 50070. The second line below is the default (port 22 is required for regular SSH connections). The third line opens all ports but ONLY for the same security group (assuming that all of your nodes in the cluster share the same security group – that will happen automatically if you use the “create more like this” option when creating instances as specified in part 1-c below). We previously had some issues with machines being hacked without that last limitation, so please don't skip this step



- c. Right click on the Master node and choose “create more like this” to create 2 more nodes with same settings. If you configure the network settings on master first, security group information will be copied.
NOTE: Hard drive size will not be copied and default to 8G unless you change it.
2. Connect to the master and set up Hadoop similarly to what you did previously. Do not attempt to repeat these steps on workers yet – you will only need to set up Hadoop once.
 - a. Configure core-site.xml, adding the **PrivateIP** (do not use public IP) of the master.

```
limitations under the License. See accompanying LICENSE file.
-->

<!-- Put site-specific property overrides in this file. -->

<configuration>

<property>
<name>fs.defaultFS</name>
<value>hdfs://172.31.7.201/</value>
</property>

</configuration>
[ec2-user@ip-172-31-7-201 ~]$ cat hadoop-2.6.4/etc/hadoop/core-site.xml
```

- b. Configure hdfs-site and set replication factor to 2.

```
<!-- Put site-specific property overrides in this file. -->

<configuration>

<property>
<name>dfs.replication</name>
<value>2</value>
</property>

</configuration>
[ec2-user@ip-172-31-9-105 ~]$
```

- c. cp hadoop-2.6.4/etc/hadoop/mapred-site.xml.template hadoop-2.6.4/etc/hadoop/mapred-site.xml and then configure mapred-site.xml

```
<!-- Put site specific property overrides in this file. -->

<configuration>

<property>
<name>mapreduce.framework.name</name>
<value>yarn</value>
</property>

</configuration>
[ec2-user@ip-172-31-9-105 ~]$ cat hadoop-2.6.4/etc/hadoop/mapred-site.xml
```

- d. Configure yarn-site.xml (once again, use PrivateIP of the master)

```
<!-- Site specific YARN configuration properties -->

<property>
<name>yarn.resourcemanager.hostname</name>
<value>172.31.7.201</value>
</property>

<property>
<name>yarn.nodemanager.aux-services</name>
<value>mapreduce_shuffle</value>
</property>

</configuration>
[ec2-user@ip-172-31-7-201 ~]$ cat hadoop-2.6.4/etc/hadoop/yarn-site.xml
```

Finally, edit the slaves file and list your 3 nodes (master and 2 workers) using Private IPs

```
[ec2-user@ip-172-31-7-201 ~]$ cat hadoop-2.6.4/etc/hadoop/slaves
172.31.7.201
172.31.5.246
...
```

Make sure that you use private IP (private DNS is also ok) for your configuration files (such as conf/masters and conf/slaves or the other 3 config files). The advantage of the Private IP is that it does not change after your instance is stopped (if you use the Public IP, the cluster would need to be reconfigured every time it is stopped). The downside of the Private IP is that it is only meaningful within the Amazon EC2 network. So all nodes in EC2 can talk to each other using

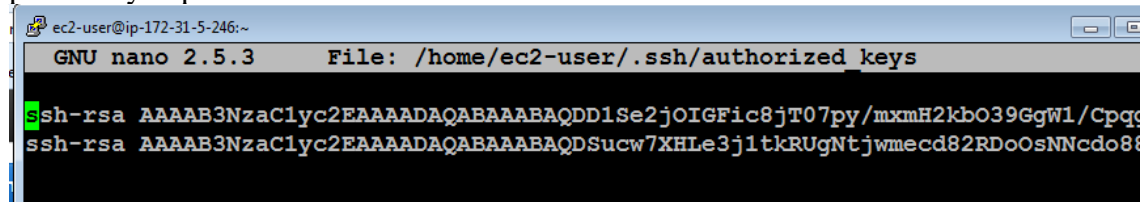
Private IP, but you cannot connect to your instance from the outside (e.g., from your laptop) because Private IP has no meaning for your laptop (since your laptop is not part of the Amazon EC2 network).

Now, we will pack up and move Hadoop to the workers. All you need to do is to generate and then copy the public key to the worker nodes to achieve passwordless access across your cluster.

1. Run `ssh-keygen -t rsa` (and enter empty values for the passphrase) on the master node. That will generate `.ssh/id_rsa` and `.ssh/id_rsa.pub` (private and public key). You now need to manually copy the `.ssh/id_rsa.pub` and append it to `~/.ssh/authorized_keys` **on each worker.**

Keep in mind that this is a single-line public key and accidentally introducing a line break would cause a mismatch.

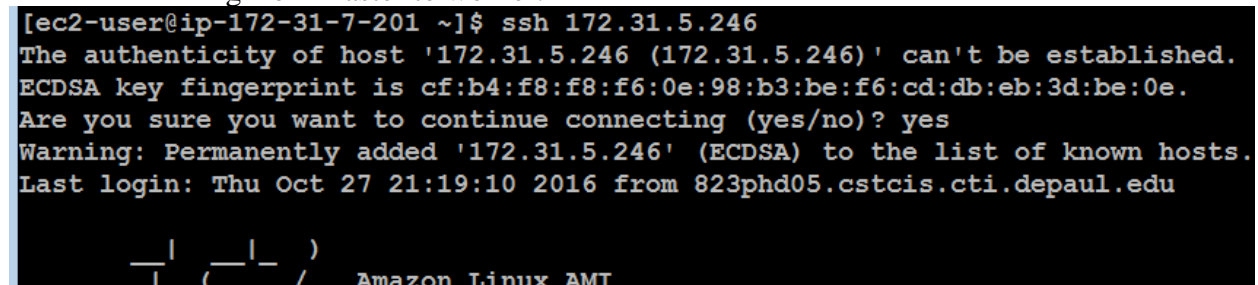
Note that the example below is NOT the master, but one of the workers (ip-172-31-5-246). The first public key is the .pem Amazon half and the 2nd public key is the master's public key copied in as one line.



```
ec2-user@ip-172-31-5-246:~  
GNU nano 2.5.3 File: /home/ec2-user/.ssh/authorized_keys  
ssh-rsa AAAAB3NzaC1yc2EAAAADAQABAAQDD1Se2jOIGFic8jT07py/mxmH2kbO39GgW1/Cpq  
ssh-rsa AAAAB3NzaC1yc2EAAAADAQABAAQDSucw7XHLe3j1tkRUGNtjwmecd82RDoOsNNcdo88
```

You can add the public key of the master to the master by running this command:
`cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys`

Make sure that you can ssh to all of the nodes from the master node (by running ssh 54.186.221.92, where the IP address is your worker node) from the master and ensuring that you were able to login. You can exit after successful ssh connection by typing exit (the command prompt will tell you which machine you are connected to, e.g., `ec2-user@ip-172-31-37-113`). Here's me ssh-ing from master to worker.



```
[ec2-user@ip-172-31-7-201 ~]$ ssh 172.31.5.246  
The authenticity of host '172.31.5.246 (172.31.5.246)' can't be established.  
ECDSA key fingerprint is cf:b4:f8:f8:f6:0e:98:b3:be:f6:cd:db:eb:3d:be:0e.  
Are you sure you want to continue connecting (yes/no)? yes  
Warning: Permanently added '172.31.5.246' (ECDSA) to the list of known hosts.  
Last login: Thu Oct 27 21:19:10 2016 from 823phd05.cstcis.cti.depaul.edu  
  
_ | _ | _ )  
_ | ( _ | _ / Amazon Linux AMI
```

Once you have verified that you can ssh from the master node to every cluster member including the master itself (ssh localhost), you are going to return to the master node (**exit** until your prompt shows the IP address of the master node) and pack the contents of the hadoop directory there. Make sure your Hadoop installation is configured correctly (because from now on, you will have 4 copies of the Hadoop directory and all changes need to be applied in 4 places).

`/usr/lib/jvm/java-1.8.0-openjdk-1.8.0.272.b10-1.amzn2.0.1.x86_64/jre/bin/java`

cd (go to root home directory, i.e. /home/ec2-user/)

(pack up the entire Hadoop directory into a single file for transfer. You can optionally compress the file with gzip)

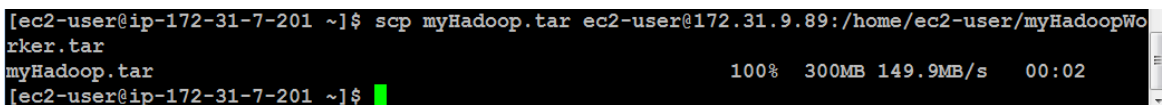
tar cvf myHadoop.tar hadoop-2.6.4

ls -al myHadoop.tar (to verify that the .tar file had been created)

Now, you need to copy the myHadoop.tar file to every non-master node in the cluster. If you had successfully setup public-private key access in the previous step, this command (for each worker node) will do that:

(copies the myHadoop.tar file from the current node to a remote node into a file called myHadoopWorker.tar. Don't forget to replace the IP address with that your worker nodes. By the way, since you are on the Amazon EC2 network, either Public or Private IP will work just fine.)

scp myHadoop.tar ec2-user@54.187.63.189:/home/ec2-user/myHadoopWorker.tar

A terminal window screenshot showing the execution of the scp command. The command is: scp myHadoop.tar ec2-user@172.31.9.89:/home/ec2-user/myHadoopWorker.tar. The output shows the file being transferred at 149.9MB/s, reaching 100% completion in 00:02. The prompt returns to the user on the master node.

```
[ec2-user@ip-172-31-7-201 ~]$ scp myHadoop.tar ec2-user@172.31.9.89:/home/ec2-user/myHadoopWorker.tar
myHadoop.tar                                100% 300MB 149.9MB/s   00:02
[ec2-user@ip-172-31-7-201 ~]$
```

Once the tar file containing your Hadoop installation from master node has been copied to each worker node, you need to login to each non-master node and unpack the .tar file.

Run the following command (on each worker node, not on the master) to untar the hadoop file. We are purposely using a different tar archive name (i.e., **myHadoopWorker.tar**), so if you get “file not found” error, that means you are running this command on the master node or have not yet successfully copied myHadoopWorker.tar file to the worker.

tar xvf myHadoopWorker.tar

Once you are done, run this on the master (nothing needs to be done on the workers to format the cluster unless you are re-formatting, in which case you'll need to delete the dfs directory).

hadoop namenode -format

Once you have successfully completed the previous steps, you should can start and use your new cluster by going to the master node and running the start-dfs.sh and start-yarn.sh scripts (you do not need to explicitly start anything on worker nodes – the master will do that for you).

You should verify that the cluster is running by pointing your browser to the link below.

http://[insert-the-public-ip-of-master]:50070/
http://3.227.208.196:50070/

Make sure that the cluster is operational (you can see the 3 nodes under Datanodes tab).

Submit a screenshot of your cluster status view.

Desktop/Capst...PDF_Search_R...PDF_Search_R...Inbox (104) - C...Mail - Craig, C...AWS Account...Workbench...Instance detai...Namenode info...+

← → ↻ Not Secure | 3.227.208.196:50070/dfshealth.html#tab-datanode

Apps Campus Connect...Mail - Craig, Chris...D2LIntelligent Informa...CSC 555 - Mining...The New York Tim...The Wall Street Jo...Chicago Tribune:...All of Statistics»

HadoopOverviewDatanodesSnapshotStartup ProgressUtilities

Datanode Information

In operation

Node	Last contact	Admin State	Capacity	Used	Non DFS Used	Remaining	Blocks	Block pool used	Failed Volumes	Version
ip-172-31-75-78.ec2.internal (172.31.75.78:50010)	0	In Service	29.99 GB	4 KB	2.22 GB	27.76 GB	0	4 KB (0%)	0	2.6.4
ip-172-31-70-33.ec2.internal (172.31.70.33:50010)	0	In Service	29.99 GB	4 KB	2.22 GB	27.76 GB	0	4 KB (0%)	0	2.6.4
ip-172-31-71-141.ec2.internal (172.31.71.141:50010)	0	In Service	29.99 GB	8 KB	2.41 GB	27.58 GB	0	8 KB (0%)	0	2.6.4

Decomissioning

Node	Last contact	Under replicated blocks	Blocks with no live replicas	Under Replicated Blocks In files under construction
------	--------------	-------------------------	------------------------------	--

Hadoop, 2014.

Legacy UI

Repeat the steps for wordcount using bioproject.xml from Assignment 1 and submit screenshots of running it.

A terminal window titled "Downloads — ec2-user@ip-172-31-71-141:~ — ssh -i ccraig13.pem ec2-user@ec2-3-236-1..." displays the following commands and output:

```
odemanager-ip-172-31-71-141.ec2.internal.out
[ec2-user@ip-172-31-71-141 ~]$ mr-jobhistory-daemon.sh start historyserver
starting historyserver, logging to /home/ec2-user/hadoop-2.6.4/logs/mapred-ec2-user-historyserv
r-ip-172-31-71-141.ec2.internal.out
[ec2-user@ip-172-31-71-141 ~]$ jps
3920 ResourceManager
4066 NodeManager
4393 JobHistoryServer
3421 NameNode
3773 SecondaryNameNode
3582 DataNode
4430 Jps
[ec2-user@ip-172-31-71-141 ~]$ wget http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/bioproject.x
1
--2021-02-11 13:21:23-- http://rasinsrv07.cstcis.cti.depaul.edu/CSC555/bioproject.xml
Resolving rasinsrv07.cstcis.cti.depaul.edu (rasinsrv07.cstcis.cti.depaul.edu)... 140.192.39.95
Connecting to rasinsrv07.cstcis.cti.depaul.edu (rasinsrv07.cstcis.cti.depaul.edu)|140.192.39.95
:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 231149003 (220M) [text/xml]
Saving to: 'bioproject.xml'

100%[=====>] 231,149,003 11.2MB/s in 20s

2021-02-11 13:21:43 (11.2 MB/s) - 'bioproject.xml' saved [231149003/231149003]

[ec2-user@ip-172-31-71-141 ~]$ hadoop fs -put bioproject.xml /data/
put: `/data/': No such file or directory
[ec2-user@ip-172-31-71-141 ~]$ hadoop fs -mkdir /data
[ec2-user@ip-172-31-71-141 ~]$ hadoop fs -put bioproject.xml /data/
[ec2-user@ip-172-31-71-141 ~]$ hadoop fs -ls /data
Found 1 items
-rw-r--r-- 2 ec2-user supergroup 231149003 2021-02-11 13:23 /data/bioproject.xml
[ec2-user@ip-172-31-71-141 ~]$
```



```

Downloads — ec2-user@ip-172-31-71-141:~ — ssh -i ccraig13.pem ec2-user@ec2-3-236-1...

File System Counters
  FILE: Number of bytes read=59605201
  FILE: Number of bytes written=86827958
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=231153309
  HDFS: Number of bytes written=20056175
  HDFS: Number of read operations=9
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=2

Job Counters
  Launched map tasks=2
  Launched reduce tasks=1
  Data-local map tasks=2
  Total time spent by all maps in occupied slots (ms)=39142
  Total time spent by all reduces in occupied slots (ms)=6355
  Total time spent by all map tasks (ms)=39142
  Total time spent by all reduce tasks (ms)=6355
  Total vcore-milliseconds taken by all map tasks=39142
  Total vcore-milliseconds taken by all reduce tasks=6355
  Total megabyte-milliseconds taken by all map tasks=40081408
  Total megabyte-milliseconds taken by all reduce tasks=6507520

Map-Reduce Framework
  Map input records=5284546
  Map output records=18562366
  Map output bytes=279356680
  Map output materialized bytes=26902454
  Input split bytes=210
  Combine input records=20053191
  Combine output records=2673165
  Reduce input groups=1040390
  Reduce shuffle bytes=26902454
  Reduce input records=1182340
  Reduce output records=1040390
  Spilled Records=3855505
  Shuffled Maps =2
  Failed Shuffles=0
  Merged Map outputs=2
  GC time elapsed (ms)=522
  CPU time spent (ms)=43600
  Physical memory (bytes) snapshot=777437184
  Virtual memory (bytes) snapshot=6418468864
  Total committed heap usage (bytes)=562561024

Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=231153099

File Output Format Counters
  Bytes Written=20056175

real    0m37.737s
user    0m3.883s
sys     0m0.354s
[ec2-user@ip-172-31-71-141 ~]$

```



```

Downloads — ec2-user@ip-172-31-71-141:~ — ssh -i ccraig13.pem ec2-user@ec2-3-236-150-206.compute-1.amazona...
[ec2-user@ip-172-31-71-141 ~]$ hadoop fs -du /data/wordcount1/
0      /data/wordcount1/_SUCCESS
20056175 /data/wordcount1/part-r-00000
[ec2-user@ip-172-31-71-141 ~]$ hadoop fs -cat /data/wordcount1/part-r-00000 | grep arctic
<I>hol<arctica</I> 28
<I>hol<arctica</I></B>. 8
<I>hol<arctica</I>; 1
<I>pale<arctica</I> 4
<I>hol<arctica</i> 1
(Antarctic 3
(Antarctica) 1
(Antarctica), 11
<Label>Antarctic 1
<Name>Antarctic 3
<Name>Antarctica 1
<Strain>Antarctic 1
<Title>Antarctic 5
Antarctic 137
Antarctic, 1
Antarctic. 2
Antarctic.</Description> 1
Antarctic.</Title> 1
Antarctic</Title> 4
Antarctica 16
Antarctica)</Title> 1
Antarctica, 9
Antarctica. 24
Antarctica.&#x0D; 3
Antarctica.</Description> 19
Antarctica</Description> 2
Antarctica</Name> 1
Antarctica</Title> 6
Palearctic 1
Project">Antarctic 1
Subarctic 11
abbr="Antarctic 1
antarctic 5
antarctica 17
antarctica</i></b>.&#x0D; 2
antarctica, 4
antarctica</Name> 10
antarctica</OrganismName> 11
antarctica</Title> 1
antarcticum 32
antarcticum</Name> 3
antarcticum</OrganismName> 3
antarcticus 31
antarcticus</i> 4
antarcticus</i></b>. 1
antarcticus). 1
antarcticus, 1
antarcticus</Name> 5
antarcticus</OrganismName> 5
arctic 21
arctica 27
arctica</I>) 2

```

Submit a short paragraph with a discussion about how the results compare (faster? slower? How much faster/slower? Due to what?)

The results are much faster with the time being 1m13.7s with one node to 37.7s with three nodes. So, the time is cut by 51%. This is due to the fact that the nodes are able to split the work amongst each other. Instead of just one node working the task, there is now 3 that are able to separate the task, and therefore, reduce time.

Part 2: Hive

- 1) Run the following query in Hive and report the time it takes to execute:

```
select lo_orderdate, sum(lo_extendedprice) as revenue
from lineorder, dwwdate
where lo_orderdate = d_datekey
    and d_year = 1996
    and lo_discount between 3 and 4
    and lo_quantity < 26
GROUP BY lo_orderdate;
```

Time taken: 26.847 seconds, Fetched: 366 row(s)

- 2) Perform the following transform operation using SELECT TRANSFORM on the dwwdate table by creating a new table. The new dwwdate table will combine d_daynuminweek, d_daynuminmonth, and d_daynuminyear into a single column in the new table. You should also eliminate of the last 3 columns (d_lastdayinmonthfl, d_holidayfl, and d_weekdayfl). The final table will have fewer columns than the original table because you merge 3 columns into 1 and remove 3 columns.

```
create table dwwdate (
  d_datekey          int,
  d_date             varchar(19),
  d_dayofweek        varchar(10),
  d_month            varchar(10),
  d_year             int,
  d_yearmonthnum     int,
  d_yearmonth        varchar(8),
  d_daynuminweek     int,
  d_daynuminmonth    int,
  d_daynuminyear     int,
  d_monthnuminyear   int,
  d_weeknuminyear    int,
  d_sellingseason     varchar(13),
  d_lastdayinweekfl  varchar(1),
  d_lastdayinmonthfl varchar(1),
  d_holidayfl        varchar(1),
  d_weekdayfl        varchar(1)
)
ROW FORMAT DELIMITED FIELDS
TERMINATED BY '|' STORED AS TEXTFILE;

LOAD DATA LOCAL INPATH '/home/ec2-user/dwwdate.tbl'
OVERWRITE INTO TABLE dwwdate;

create table dwwdate_transform (
```

Christian Craig Csc 555 project phase1

```
d_datekey      int,  
d_date         varchar(19),  
d_dayofweek    varchar(10),  
d_month        varchar(10),  
d_year         int,  
d_yearmonthnum int,  
d_yearmonth    varchar(8),  
d_wmy          int,  
d_monthnuminyear int,  
d_weeknuminyear int,  
d_sellingseason varchar(13),  
d_lastdayinweekfl varchar(1)  
  
)
```

```
ROW FORMAT DELIMITED FIELDS  
TERMINATED BY '\t' STORED AS TEXTFILE;
```

```
INSERT OVERWRITE TABLE dwdate_transform
```

```
SELECT
```

```
TRANSFORM(d_datekey,d_date,d_dayofweek,d_month,d_year,d_yearmonthnum,d_yearmonth,d_dayn  
uminweek,d_daynuminmonth,d_daynuminyear,d_monthnuminyear,d_weeknuminyear,d_sellingseason,  
d_lastdayinweekfl,d_lastdayinmonthfl,d_holidayfl,d_weekdayfl) USING 'python wmy_mapper.py'
```

```
AS
```

```
(d_datekey,d_date,d_dayofweek,d_month,d_year,d_yearmonthnum,d_yearmonth,d_wmy,d_monthnumi  
nyear,d_weeknuminyear,d_sellingseason,d_lastdayinweekfl) FROM dwdate;
```

Christian Craig Csc 555 project phase1

```

GNU nano 2.9.8                                why_sapper.py
//user/bin/python
import sys

for line in sys.stdin:
    line = line.strip().split('\t')
    #print(line[7:10])
    w = line[7]
    m = line[8]
    y = line[9]
    new_wmy = ''.join([w,m,y])
    line[7] = new_wmy

    print '\t'.join([line[0],line[1],line[2],line[3],line[4],line[5],line[6],line[7],line[10],line[11],line[12]])

```

hristmas	0													
19981218	0	December 18, 1998	Saturday	December	1998	199812	Dec1998	718352						
2 51	1	Christmas												
19981219	0	December 19, 1998	Sunday	December	1998	199812	Dec1998	119353	12	51				
hristmas	0													
19981220	0	December 20, 1998	Monday	December	1998	199812	Dec1998	220354	12	51				
hristmas	0													
19981221	0	December 21, 1998	Tuesday	December	1998	199812	Dec1998	321355	12	51				
hristmas	0													
19981222	0	December 22, 1998	Wednesday	December	1998	199812	Dec1998	422356						
2 51	0	Christmas												
19981223	0	December 23, 1998	Thursday	December	1998	199812	Dec1998	523357						
2 52	0	Christmas												
19981224	0	December 24, 1998	Friday	December	1998	199812	Dec1998	624358	12	52				
hristmas	0													
19981225	0	December 25, 1998	Saturday	December	1998	199812	Dec1998	725359						
2 52	1	Christmas												
19981226	0	December 26, 1998	Sunday	December	1998	199812	Dec1998	126360	12	52				
hristmas	0													
19981227	0	December 27, 1998	Monday	December	1998	199812	Dec1998	227361	12	52				
hristmas	0													
19981228	0	December 28, 1998	Tuesday	December	1998	199812	Dec1998	328362	12	52				
hristmas	0													
19981229	0	December 29, 1998	Wednesday	December	1998	199812	Dec1998	429363						
2 52	0	Christmas												
19981230	0	December 30, 1998	Thursday	December	1998	199812	Dec1998	530364						
2 53	0	Christmas												

Time taken: 0.558 seconds, Fetched: 2556 row(s)

hive>

Part 3: Pig

Convert and load the data into Pig, implementing and timing the following queries: One easy way to time Pig is as follows: put your sequence of pig commands, including LOAD, into a text file and then run, from command line in pig directory (e.g., [ec2-user@ip-172-31-6-39 pig-0.15.0]\$), **bin/pig -f pig_script.pig** (which will report how long the pig script took to run).

```
SELECT lo_discount, COUNT(lo_extendedprice)
```

Christian Craig Csc 555 project phase1

```
FROM lineorder
GROUP BY lo_discount;
```

bin/pig

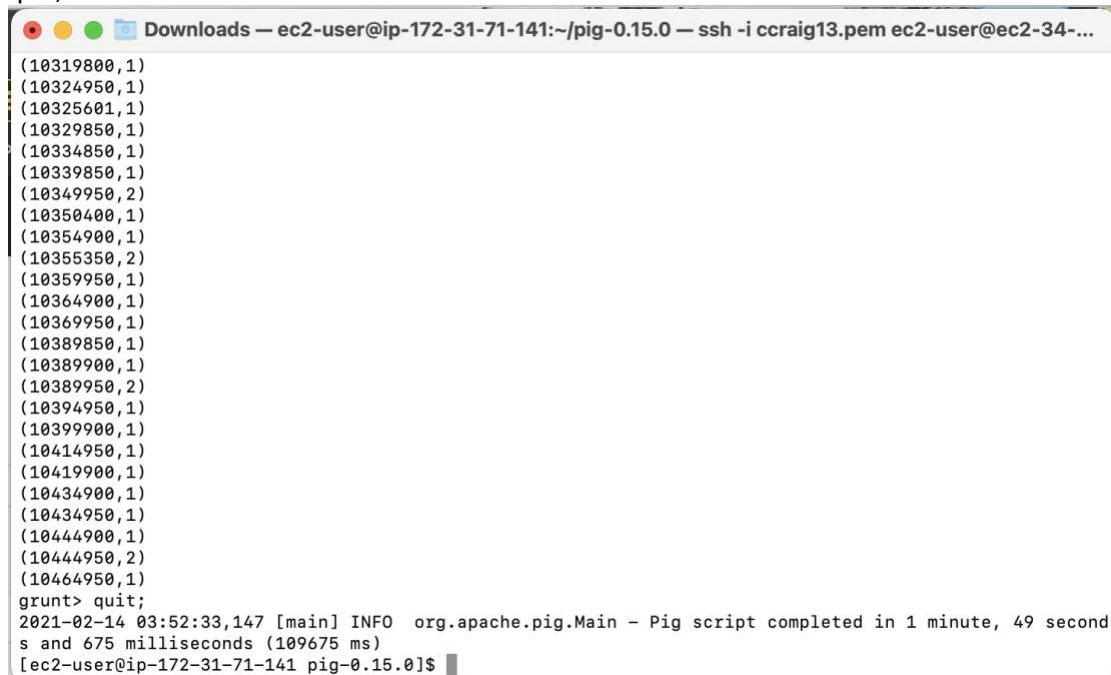
```
lineorder_tbl = LOAD '/user/ec2-user/lineorder.tbl' USING PigStorage('|') AS(lo_orderkey:int,
lo_linenummer:int, lo_custkey: int, lo_partkey:int, lo_suppkey:int,lo_orderdate:
int,lo_orderpriority:chararray,lo_shippriority:chararray,lo_quantity:int,lo_extendedprice:int,lo_ordertotalpri
ce:int,lo_discount: int,lo_revenue: int,lo_supplycost:int,lo_tax: int,lo_commitdate: int,
lo_shipmode:chararray);
```

```
lo_group1 = group lineorder_tbl by $12;
```

```
lo_count1 = foreach lo_group1 generate group as disc, COUNT(lineorder_tbl.lo_extendedprice);
```

```
dump lo_count1;
```

```
quit;
```



```
(10319800,1)
(10324950,1)
(10325601,1)
(10329850,1)
(10334850,1)
(10339850,1)
(10349950,2)
(10350400,1)
(10354900,1)
(10355350,2)
(10359950,1)
(10364900,1)
(10369950,1)
(10389850,1)
(10389900,1)
(10389950,2)
(10394950,1)
(10399900,1)
(10414950,1)
(10419900,1)
(10434900,1)
(10434950,1)
(10444900,1)
(10444950,2)
(10464950,1)
grunt> quit;
2021-02-14 03:52:33,147 [main] INFO org.apache.pig.Main - Pig script completed in 1 minute, 49 second
s and 675 milliseconds (109675 ms)
[ec2-user@ip-172-31-71-141 pig-0.15.0]$
```

1m49s

```
SELECT lo_quantity, SUM(lo_revenue)
FROM lineorder
WHERE lo_discount > 6 AND lo_quantity > 33
```

```
GROUP BY lo_quantity;
```

```
lineorder_tbl = LOAD '/user/ec2-user/lineorder.tbl' USING PigStorage('|') AS(lo_orderkey:int,  
lo_linenummer:int, lo_custkey:int, lo_partkey:int,  
lo_suppkey:int,lo_orderdate:int,lo_orderpriority:chararray,lo_shippriority:chararray,lo_quantity:int,lo_exte  
ndedprice:int,lo_ordertotalprice:int,lo_discount:  
int,lo_revenue:int,lo_supplycost:int,lo_tax:int,lo_commitdate:int, lo_shipmode:chararray);
```

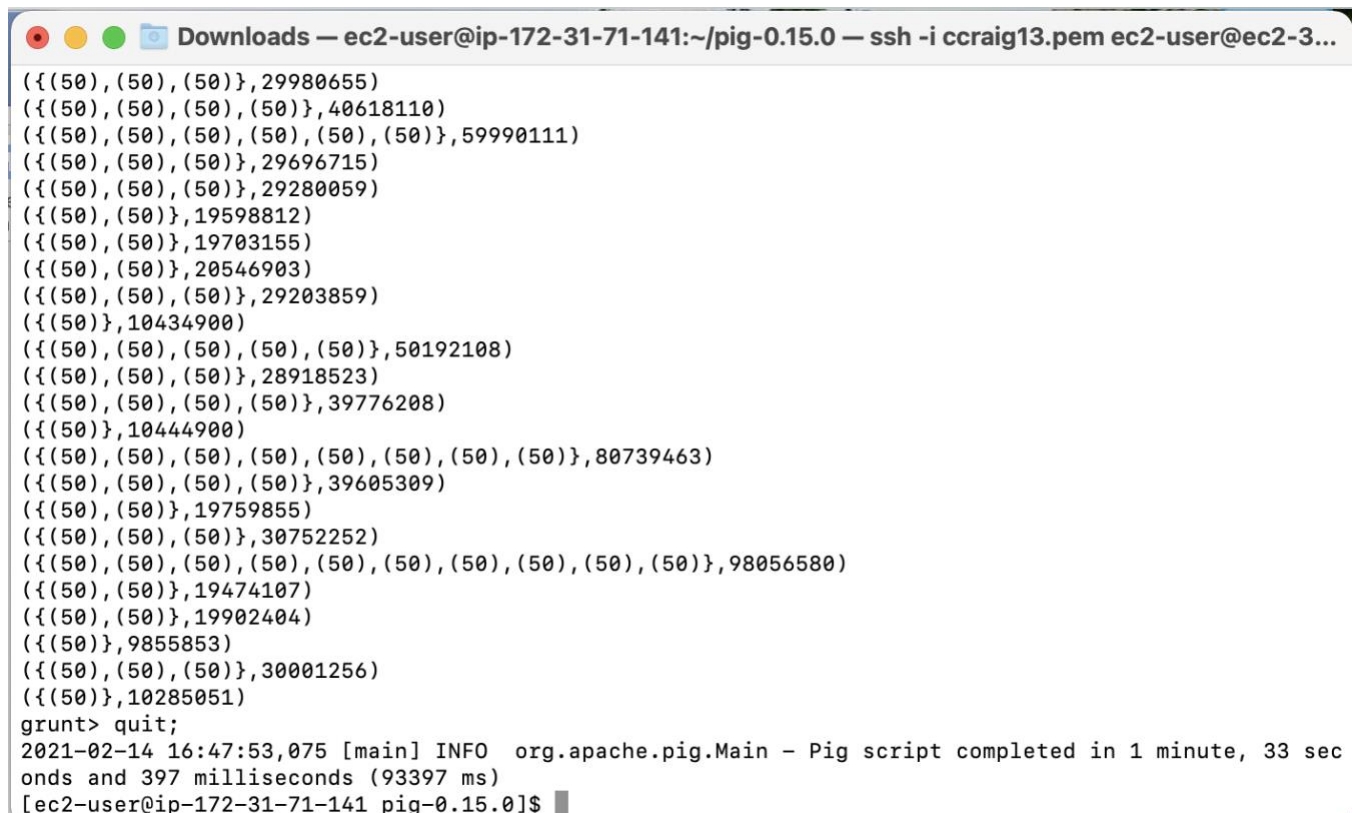
```
lo_filter = filter lineorder_tbl by $12 > 6 AND $9 > 33;
```

```
lo_group2 = group lo_filter by $9;
```

```
lo_sum = foreach lo_group2 generate lo_filter.lo_quantity, SUM(lo_filter.lo_revenue);
```

```
dump lo_sum;
```

```
quit;
```



The screenshot shows a terminal window titled "Downloads — ec2-user@ip-172-31-71-141:~/pig-0.15.0 — ssh -i ccraig13.pem ec2-user@ec2-3...". The terminal displays the output of a Pig script, which consists of 20 rows of data. Each row is a tuple of values enclosed in curly braces and separated by commas. The values include integers and long integers. After the data output, the prompt "grunt> quit;" is shown, followed by a status message: "2021-02-14 16:47:53,075 [main] INFO org.apache.pig.Main - Pig script completed in 1 minute, 33 seconds and 397 milliseconds (93397 ms)". The final prompt is "[ec2-user@ip-172-31-71-141 pig-0.15.0]\$".

```
{(50),(50),(50)},29980655)  
{(50),(50),(50),(50)},40618110)  
{(50),(50),(50),(50),(50)},59990111)  
{(50),(50),(50)},29696715)  
{(50),(50),(50)},29280059)  
{(50),(50)},19598812)  
{(50),(50)},19703155)  
{(50),(50)},20546903)  
{(50),(50),(50)},29203859)  
{(50)},10434900)  
{(50),(50),(50),(50),(50)},50192108)  
{(50),(50),(50)},28918523)  
{(50),(50),(50),(50)},39776208)  
{(50)},10444900)  
{(50),(50),(50),(50),(50),(50),(50),(50)},80739463)  
{(50),(50),(50),(50)},39605309)  
{(50),(50)},19759855)  
{(50),(50),(50)},30752252)  
{(50),(50),(50),(50),(50),(50),(50),(50),(50),(50)},98056580)  
{(50),(50)},19474107)  
{(50),(50)},19902404)  
{(50)},9855853)  
{(50),(50),(50)},30001256)  
{(50)},10285051)  
grunt> quit;  
2021-02-14 16:47:53,075 [main] INFO org.apache.pig.Main - Pig script completed in 1 minute, 33 seconds and 397 milliseconds (93397 ms)  
[ec2-user@ip-172-31-71-141 pig-0.15.0]$
```


1m33s

```
Part_tbl = LOAD '/user/ec2-user/part.tbl' USING PigStorage(',') AS (p_partkey:int, p_name:chararray,  
p_mfgr:chararray, p_category:chararray, p_brand1:chararray, p_color:chararray, p_type:chararray,  
p_size:int, p_container:chararray);
```

```
supplier_tbl = LOAD '/user/ec2-user/supplier.tbl' USING PigStorage('|') AS(s_suppkey:int,  
s_name:chararray, s_address:chararray, s_city:chararray,s_nation:chararray, s_region:chararray,  
s_phone:chararray);
```

```
customer_tbl = LOAD '/user/ec2-user/customer.tbl' USING PigStorage('|') AS(c_custkey:int,  
c_name:chararray, c_address:chararray, c_city:chararray,c_nation:chararray, c_region:chararray,  
c_phone:chararray, c_mktsegment:chararray);
```

```
dwdate_tbl = LOAD '/user/ec2-user/dwdate.tbl' USING PigStorage('|') AS(d_datekey:int,  
d_date:chararray, d_dayofweek:chararray, d_year:int,  
d_yearmonthnum:int,d_yearmonth:chararray,d_daynuminweek:int,d_daynuminmonth:int,d_daynuminye  
ar:int,d_monthnuminyear:int,d_weeknuminyear:int,d_sellingseason:chararray,d_lastdayinweekfl:chararr  
ay,d_lastdayinmonthfl: chararray,d_holidayfl:chararray,d_weekdayfl:chararray);
```

```
lineorder_tbl = LOAD '/user/ec2-user/lineorder.tbl' USING PigStorage('|') AS(lo_orderkey:int,  
lo_linenumber:int, lo_custkey: int, lo_partkey:int,  
lo_suppkey:int,lo_orderdate:int,lo_orderpriority:chararray,lo_shippriority:chararray,lo_quantity:int,lo_exte  
ndedprice:int,lo_ordertotalprice:int,lo_discount:  
int,lo_revenue:int,lo_supplycost:int,lo_tax:int,lo_commitdate:int, lo_shipmode:chararray);
```

Part 4: Hadoop Streaming

Implement, run and time the following query using Hadoop streaming with python. STDDEV refers to standard deviation.

```
SELECT lo_shipmode, STDDEV(lo_tax)
FROM lineorder
WHERE lo_quantity BETWEEN 15 AND 18
GROUP BY lo_shipmode;
```

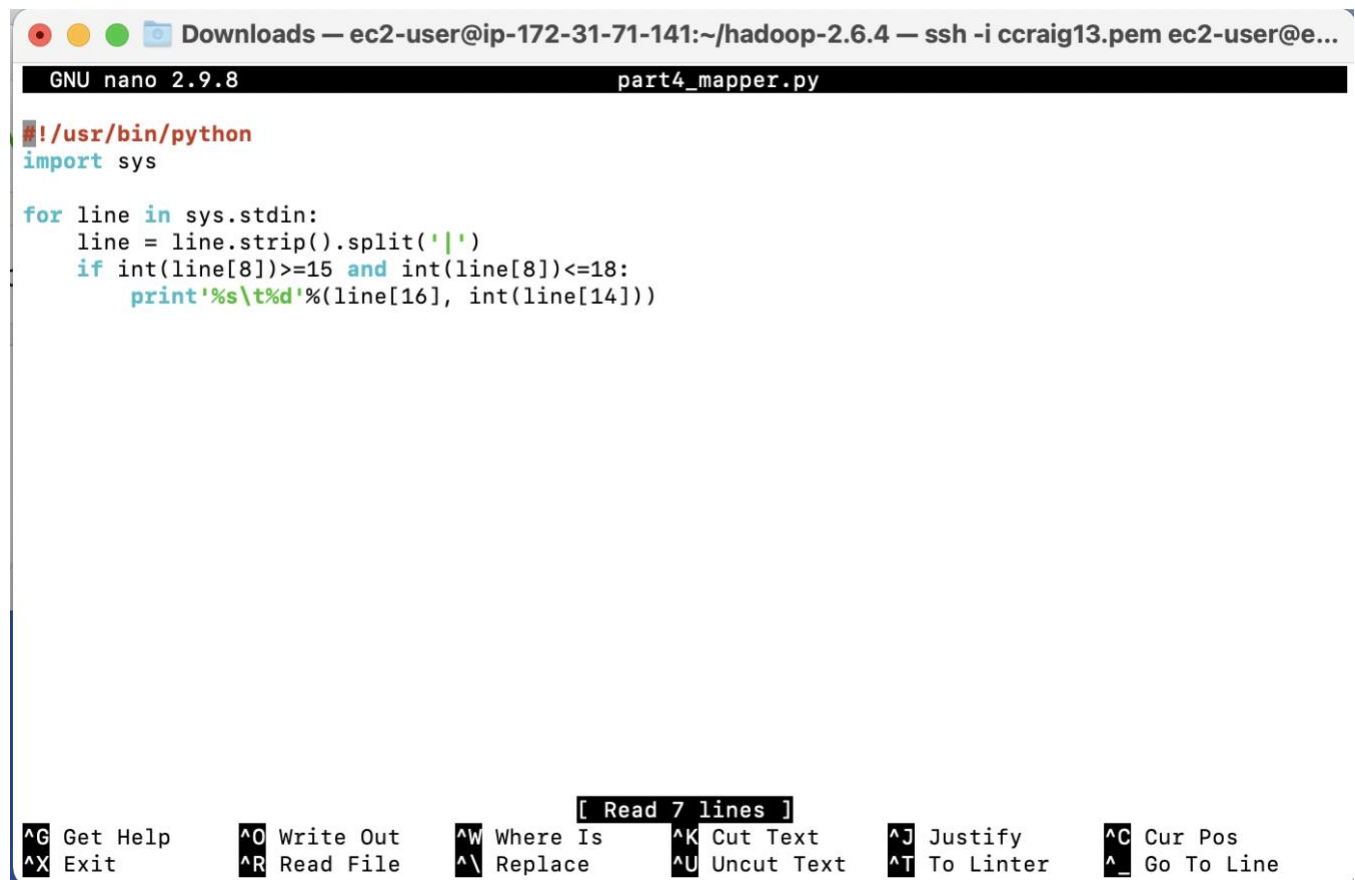
```
create table lineorder (
  lo_orderkey      int,
  lo_linenumbers   int,
  lo_custkey       int,
  lo_partkey       int,
  lo_suppkey       int,
  lo_orderdate     int,
  lo_orderpriority varchar(15),
  lo_shippriority  varchar(1),
  lo_quantity      int,
  lo_extendedprice int,
  lo_ordertotalprice int,
  lo_discount      int,
  lo_revenue       int,
  lo_supplycost    int,
  lo_tax           int,
  lo_commitdate    int,
  lo_shipmode      varchar(10)
)
ROW FORMAT DELIMITED FIELDS
TERMINATED BY '|' STORED AS TEXTFILE;
```

Christian Craig Csc 555 project phase1

LOAD DATA LOCAL INPATH '/home/ec2-user/lineorder.tbl'

OVERWRITE INTO TABLE lineorder;

-rw-r--r-- 2 ec2-user supergroup 594313001 2021-02-13 22:19 lineorder.tbl



The screenshot shows a terminal window with a title bar indicating the user is connected via SSH to an EC2 instance. The terminal is running the GNU nano 2.9.8 text editor, editing a file named `part4_mapper.py`. The script is a Python program that reads from standard input, processes each line by stripping and splitting it, and then prints specific fields based on a range check. The bottom of the screen displays a row of keyboard shortcuts for nano editor functions.

```
GNU nano 2.9.8 part4_mapper.py

#!/usr/bin/python
import sys

for line in sys.stdin:
    line = line.strip().split('|')
    if int(line[8])>=15 and int(line[8])<=18:
        print '%s\t%d'%(line[16], int(line[14]))
```

[Read 7 lines]

^G Get Help	^O Write Out	^W Where Is	^K Cut Text	^J Justify	^C Cur Pos
^X Exit	^R Read File	^_\ Replace	^U Uncut Text	^T To Linter	^_ Go To Line

NOTE: You may implement this part in Java if you p



The screenshot shows a terminal window with a title bar indicating an SSH connection to an EC2 instance. The terminal is running the GNU nano 2.9.8 text editor, editing a file named `part4_reducer.py`. The Python script is a MapReduce reducer that processes input lines, splits them by tab, and calculates the standard deviation of values for each key. The script uses `sys` for standard input and `statistics` for standard deviation calculations. The bottom of the terminal displays a row of nano editor shortcuts.

```
GNU nano 2.9.8 part4_reducer.py

#!/usr/bin/python

import sys
import statistics as st
currentKey = None
key= None
vals = []
val =None

# input comes from STDIN
for line in sys.stdin:

    split = line.strip().split('\t')
    #get key
    key = split[0]
    #value = '\t'.join(split[1:])
    val = split[1]
    if currentKey == key: # Same key
        vals.append(int(val))
    else:
        if currentKey: #get ouput
            sdev = st.stdev(vals)
            print currentKey, '\t', sdev

            currentKey = key
            vals = []
            val = split[1]
            vals.append(int(val))

if len(vals)>0:
    sdev = st.stdev(vals)
    print '%s\t%d'%(key, sdev)
```

^G Get Help ^O Write Out ^W Where Is ^K Cut Text ^J Justify ^C Cur Pos
^X Exit ^R Read File ^\ Replace ^U Uncut Text ^T To Linter ^ Go To Line

refer.

```
Downloads — ec2-user@ip-172-31-71-141:~/hadoop-2.6.4 — ssh -i ccraig13.pem ec2-user@e...  
WRONG_REDUCE=0  
File Input Format Counters  
  Bytes Read=594329385  
File Output Format Counters  
  Bytes Written=128  
21/02/15 05:17:04 INFO streaming.StreamJob: Output directory: /data/output6  
[[ec2-user@ip-172-31-71-141 hadoop-2.6.4]$ hadoop fs /user/ec2-user/data/output6  
/user/ec2-user/data/output6: Unknown command  
[[ec2-user@ip-172-31-71-141 hadoop-2.6.4]$ hadoop fs /ec2-user/data/output6  
/ec2-user/data/output6: Unknown command  
[[ec2-user@ip-172-31-71-141 hadoop-2.6.4]$ hadoop fs /data/output6  
/data/output6: Unknown command  
[[ec2-user@ip-172-31-71-141 hadoop-2.6.4]$ hadoop fs - cat /data/output6  
cat: Unknown command  
Did you mean -cat? This command begins with a dash.  
[[ec2-user@ip-172-31-71-141 hadoop-2.6.4]$ hadoop fs -cat /data/output6  
cat: `/data/output6': Is a directory  
[[ec2-user@ip-172-31-71-141 hadoop-2.6.4]$ hadoop fs -ls /data/output6  
Found 2 items  
-rw-r--r--  2 ec2-user supergroup          0 2021-02-15 05:17 /data/output6/_SUCCESS  
-rw-r--r--  2 ec2-user supergroup       128 2021-02-15 05:17 /data/output6/part-00000  
[[ec2-user@ip-172-31-71-141 hadoop-2.6.4]$ hadoop fs -cat /data/output6/part-00000  
AIR      2.58531913756  
FOB      2.58050450149  
MAIL     2.57629684481  
RAIL     2.57205496503  
REG AIR   2.57882092561  
SHIP     2.5793514077  
TRUCK    2  
[[ec2-user@ip-172-31-71-141 hadoop-2.6.4]$
```

Submit a single document containing your written answers. Be sure that this document contains your name and “CSC 555 Project Phase 1” at the top.